

# Query Planning and Optimization

Dr. Qichen Wang

EPFL

2025.5

# Self-introduction

---

- Dr. Qichen Wang
  - PhD from Hong Kong University of Science and Technology, 2022
  - Research Assistant Professor, Hong Kong Baptist University 2022-2024
  - Postdoc, EPFL, 2024-now
- Teaching experiences:
  - Lecturer: Cloud Computing, Hong Kong Baptist University
  - TA: Big Data Technology, Combinatorial Optimization, HKUST
- Teaching interests:
  - Databases, Cloud Computing, Big Data Technology, Algorithms, Data Structures
  - Other BS/MS level CS courses

# Prerequisite

---

- Fundamental relational concepts: tables, tuples, columns, primary and foreign keys
- Relational algebra
- Basic concepts of writing SQL queries, **SELECT**, **FROM**, **WHERE**, different types of joins, and subqueries
- Big-O analysis for algorithmic cost

# Demo Database

- Student(sid, name, state), Course(cid, title), Enrolled(sid, cid, grade)

sid	name	state
1	Alice	CA
2	Bob	NY
3	Charlie	CA
4	Diana	TX
5	Eve	CA
6	Frank	TX
7	Grace	NY

cid	title
101	Database Systems
102	Operating Systems
103	Algorithms
104	Computer Networks

sid	cid	grade
1	101	A
1	103	B
2	101	B
2	102	A
3	101	A
3	102	B
3	103	A
3	104	A
4	103	C
5	101	B
5	102	A
6	101	A
7	104	A
8	101	A

Download the demo database



<https://qichen-wang.github.io/files/demo.sql>

To load it:

For DuckDB:

```
.read /path/to/demo.sql
```

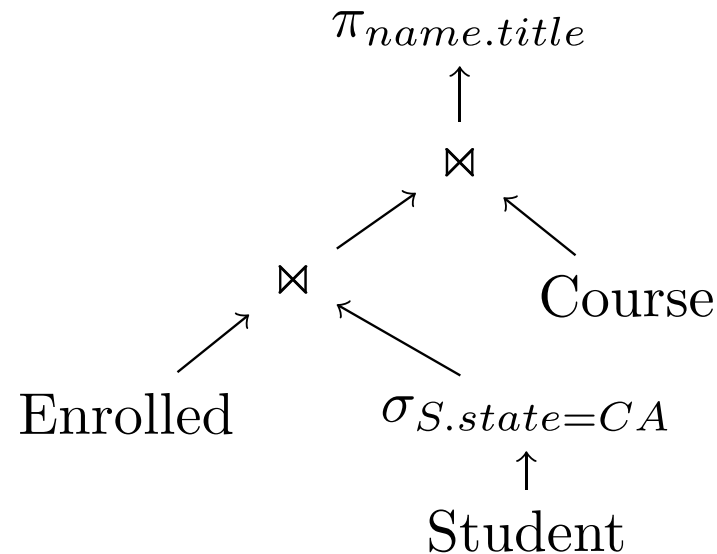
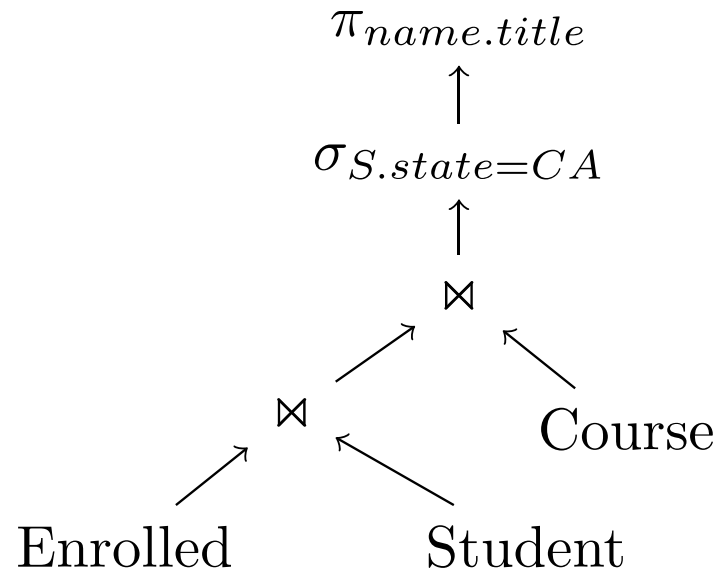
For PostgreSQL:

```
\i /path/to/demo.sql
```

# SQL: A declarative language

- When writing SQL queries, we only express our high-level ideas.
- There can be different ways of evaluating the query.
  - “Listing all students from CA and the courses they have enrolled in.”

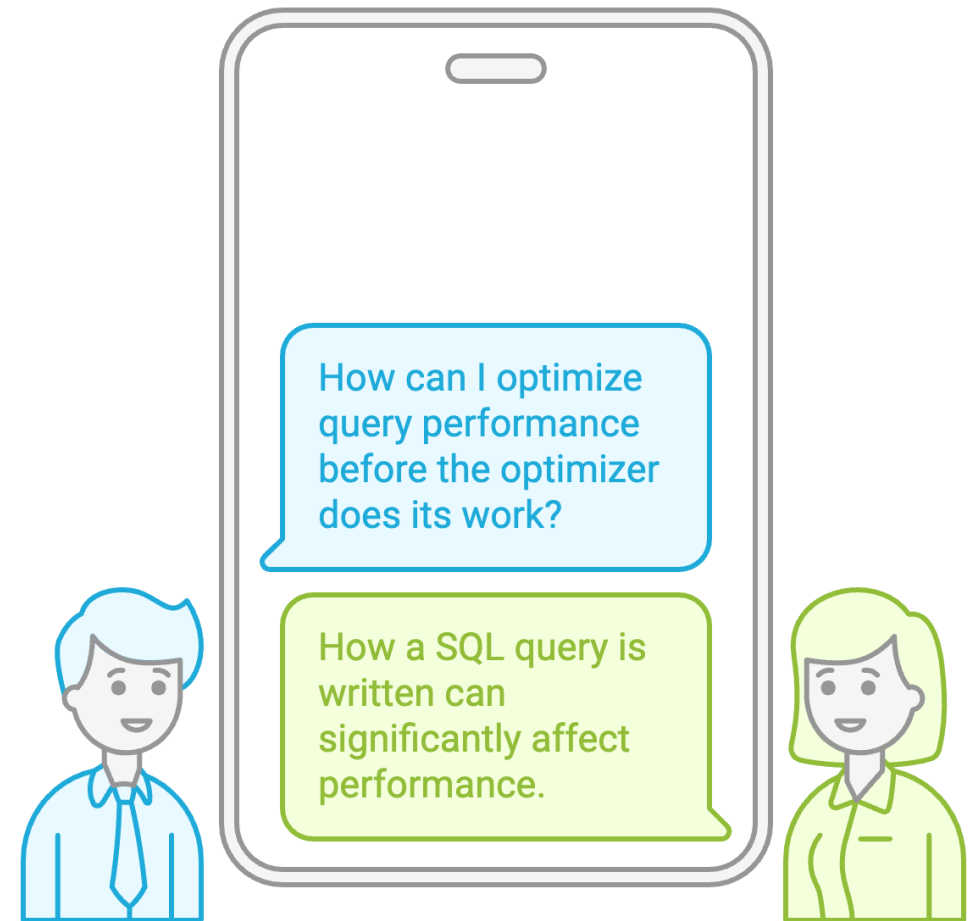
```
SELECT name, title
FROM Student s, Course c, Enrolled e
WHERE s.sid = e.sid
AND c.cid = e.cid
AND s.state = 'CA';
```



Before Optimization

# The first step of optimization

- Ideally, the optimizer should do everything for you.
  - But that is not the case for current database systems.



# An example:

Student(sid, name, state), Course(cid, title), Enrolled(sid, cid, grade)

- Suppose you want to find the students who have enrolled in all courses
- What will you do?
- 'For all' is hard to represent in SQL
- A direct translation: **Find the students for whom there are no course they have not enrolled in.**

```
SELECT sid
FROM Student s
WHERE NOT EXISTS (
  SELECT * FROM Course c
  WHERE NOT EXISTS (
    SELECT * FROM Enrolled e
    WHERE s.sid = e.sid AND c.cid = e.cid
  ));
```

It takes  $O(n^2)$  time  
Loop over all students and courses  
and check the Enrolled table for every  
possible combination.



# How to do better?

---

Student(sid, name, state), Course(cid, title), Enrolled(sid, cid, grade)

- Suppose you want to find the students who have enrolled in all courses
- Another possible way: **Find the students whose enrolled course count matches the total number of courses in the Course table.**

```
SELECT sid  
FROM Enrolled e  
GROUP BY sid  
HAVING count(*) = (SELECT count(*) FROM Course c);
```

Can be done in linear time  $O(n)$

- **Writing a good SQL can reduce the complexity at the beginning.**

# Some good practices you should know

---

## ■ Rule 1: Select Only Necessary Columns

- To avoid select \* queries.
- It is hard to find a query requiring every table column.
- For some databases, data is stored in columnar format.
- Selecting only required columns can significantly reduce the I/O cost.

# Some good practices you should know

---

- **Rule 2: Remove redundant filter conditions and avoid functions in filter conditions**
  - For example, having both "data >= 2025-01-01 and data <= 2025-12-31" and "YEAR(date) = 2025"
  - YEAR(date) = 2025 is redundant
  - Also, YEAR(date) = 2025 is not index-friendly; databases usually have indices on the range queries, but not for functions.

# Some good practices you should know

## ■ Rule 3: Replace IN with EXISTS

- For some databases, the EXISTS clause often offers better performance.

```
SELECT name
FROM Student
WHERE state = 'CA'
  AND sid IN ( SELECT E.sid
               FROM Enrolled e, Course c
               WHERE e.cid = c.cid
               AND c.title = 'Database Systems'
               AND e.grade = 'A');
```

```
SELECT name
FROM Student s
WHERE EXISTS (SELECT 1
              FROM Enrolled e, Course c
              WHERE e.cid = c.cid
              AND s.sid = e.sid
              AND c.title = 'Database Systems'
              AND e.grade = 'A')
AND state = 'CA';
```

- Some databases can optimize that for you (e.g., DuckDB) while some cannot (e.g., PostgreSQL)
- Always use EXISTS if the right-hand side is a subquery.

# Some good practices you should know

## ■ Rule 4: Replace unnecessary joins with semi-joins (EXISTS)

- Some join queries can be replaced with a semi-join if the output attributes are only located in one of the two relations.

```
SELECT DISTINCT S.name
FROM Student s, Enrolled e, Course c
WHERE S.state = 'CA'
AND C.title = 'Database Systems'
AND E.grade = 'A'
AND s.sid = e.sid AND c.cid = e.cid;
```

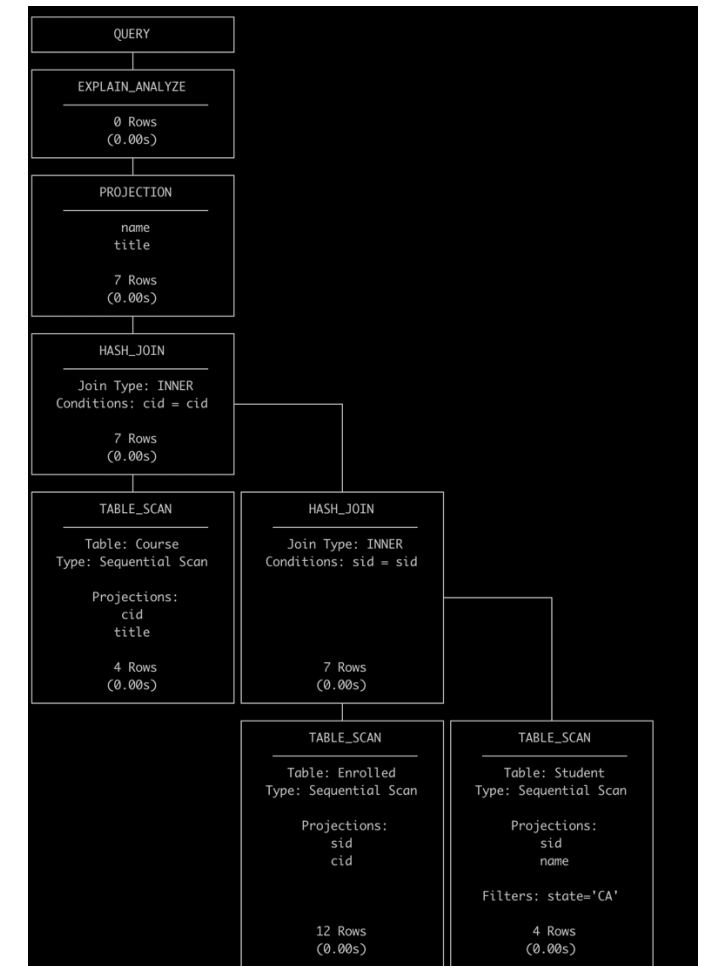
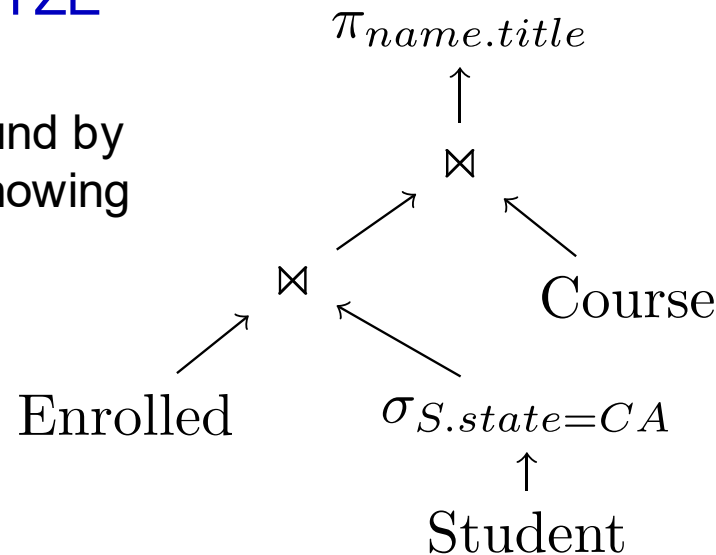
- Avoid costly full join computation.

```
SELECT name
FROM Student s
WHERE EXISTS (SELECT 1
              FROM Enrolled e
              WHERE s.sid = e.sid
              AND e.grade = 'A'
              AND EXISTS (SELECT 1
                          FROM Course c
                          WHERE c.cid = e.cid
                          AND c.title = 'Database Systems'))
AND S.state = 'CA';
```

# Viewing Query Evaluation Plans

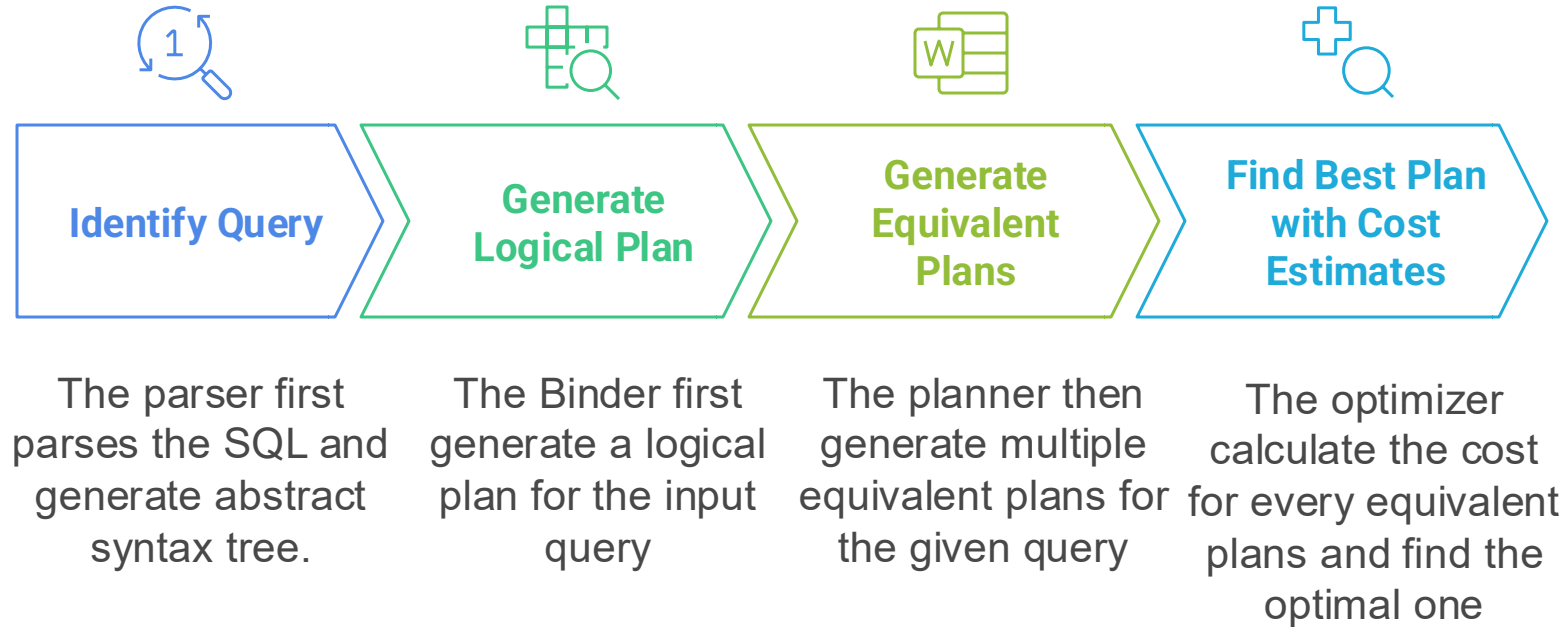
- Most databases support '**EXPLAIN** <query>' to display the query execution plan.
  - Display plan chosen by query optimizer, along with cost estimation
- Some databases (e.g., PostgreSQL, DuckDB) support '**EXPLAIN ANALYZE** <query>'
  - Shows actual runtime statistics found by running the query, in addition to showing the plan

```
EXPLAIN ANALYZE SELECT name, title
FROM Student s, Course c, Enrolled e
WHERE s.sid = e.sid
AND c.cid = e.cid
AND s.state = 'CA';
```



# Logical Plans and Rule-based Optimization

# Logical Query Optimization



- The logical plan corresponds to a relational algebra expression.
- We need to find the equivalent relational algebra expressions to find equivalent plans.



# Transformation of Relational Expressions

---

- Two relational algebra expressions are said to be equivalent if the two expressions generate the same set of tuples on every legal database instance.
  - Note: order of tuples is irrelevant
- An **equivalence rule** says that expressions of two forms are equivalent.
  - Can replace the expression of the first form by the second, or vice versa
- It is actually hard to find all possible equivalent expressions
  - NP-hard problem
- **Practically:** Choose from a subset of all possible plans

# Equivalence Rules

1. Conjunctive selection operations can be deconstructed into a sequence of individual selections.

$$\sigma_{\theta_1 \wedge \theta_2}(E) \equiv \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

2. Selection operations are commutative.

$$\sigma_{\theta_1}(\sigma_{\theta_2}(E)) \equiv \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

3. Only the last in a sequence of projection operations is needed, the others can be omitted.

$$\pi_{L_1}(\pi_{L_2}(\cdots(\pi_{L_n}(E)))) \equiv \pi_{L_1}(E)$$

where  $L_1 \subseteq L_2 \subseteq \cdots \subseteq L_n$

4. Join are commutative

$$E_1 \bowtie E_2 \equiv E_2 \bowtie E_1$$

5. Natural join are associative

$$(E_1 \bowtie E_2) \bowtie E_3 \equiv E_1 \bowtie (E_2 \bowtie E_3)$$

# Predicate Pushdown

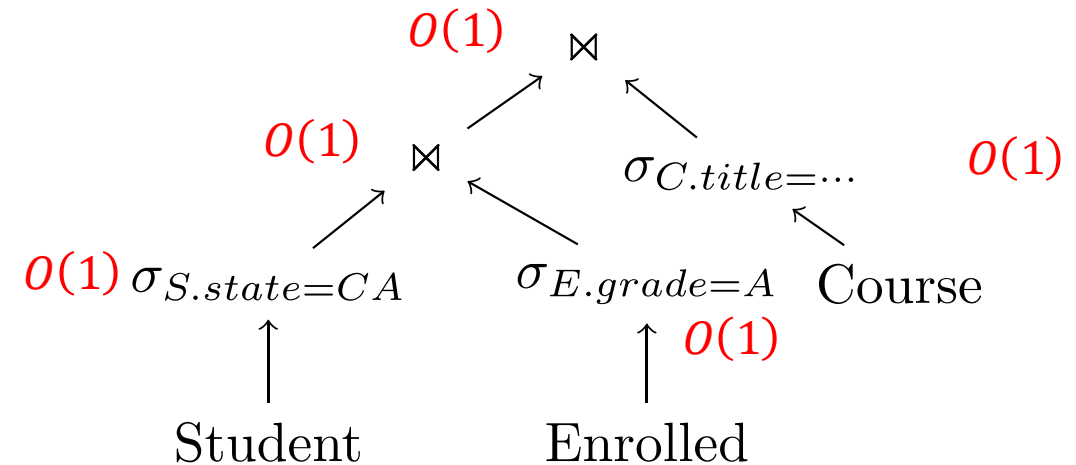

6. The selection operation can be distributed over the join operations if all the attributes in  $\theta$  involve only the attributes of one of the expressions ( $E_1$ ) being joined.

$$\sigma_{\theta}(E_1 \bowtie E_2) \equiv (\sigma_{\theta}(E_1)) \bowtie E_2$$

$$\sigma_{S.state=CA \wedge C.title=Database\ Systems \wedge E.grade=A}(s \bowtie e \bowtie c)$$

$$(\sigma_{S.state=CA}(s)) \bowtie (\sigma_{E.grade=A}(e)) \bowtie (\sigma_{C.title=Database\ System}(c))$$

```
SELECT DISTINCT S.name
FROM Student s, Enrolled e, Course c
WHERE S.state = 'CA'
AND C.title = 'Database Systems'
AND E.grade = 'A'
AND s.sid = e.sid AND c.cid = e.cid;
```



# Projection Pushdown

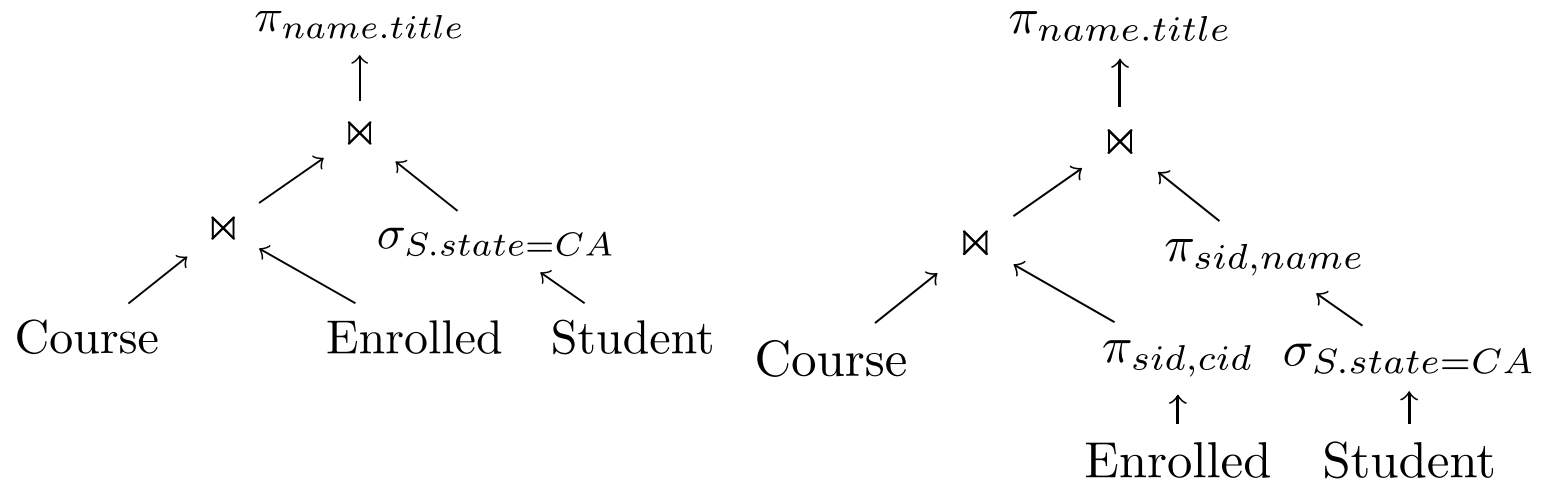
7. The projection operation distributes over the join operation as follows:

Assume  $L_1/L_2$  only involves attributes from  $E_1/E_2$ ,  $L_3$  are the set of join attributes:

$$\pi_{L_1 \cup L_2}(E_1 \bowtie E_2) \equiv \pi_{L_1 \cup L_2}(\pi_{L_1 \cup L_3}(E_1) \bowtie \pi_{L_2 \cup L_3}(E_2))$$

i.e., we first project all attributes in  $E_1/E_2$  that are either not in the final output attributes, or the join attributes. After calculating the join, we remove all the non-output join attributes ( $\pi_{L_1 \cup L_2}$ )

```
SELECT name, title
FROM Student s, Course c, Enrolled e
WHERE s.sid = e.sid
AND c.cid = e.cid
AND s.state = 'CA';
```



# Take-home exercise

---

- Can you use projection pushdown and the following rule

$$\left(E_1 \bowtie (\pi_{L_3} E_2)\right) \equiv E_1 \bowtie E_2$$

to find an equivalent rule for replacing joins with semi-joins?

$$\pi_{L_1}(E_1 \bowtie E_2) \equiv (\pi_{L_1} E_1) \ltimes E_2$$

- $L_1$  only involves attributes from  $E_1$
- $L_3$  are the join attributes between  $E_1$  and  $E_2$

# Heuristic Optimizations

---

- There are more rules (even rules that have not been discovered yet).
- These techniques **do not need** to examine data.
  - Predicate pushdown
  - Projection pushdown
- Idea: drop unused data as much as possible and as early as possible without affecting the efficiency
- Provide a much better starting point for the next stage of optimization.

# Cost-based Optimization

# Cost-based Query Optimization

---

- The efficiency of a query plan depends on multiple factors:
  - CPU time
  - I/O operations
  - Memory usage
  - Cache misses
- Cost Model: a weighted formula that combines all these factors:
$$c_1(\text{CPU Ops}) + c_2(\text{I/O Ops}) + \dots$$
  - The constants  $c_1, c_2, \dots$  depend heavily on hardware
  - They are determined by the database system.
  - The formula can be simpler or more complicated.
- Also, heavily depends on the output size of each operator, which determine the number of CPU and I/O operations



# Cost Estimation

---

- Need statistics of input relations.
  - E.g., number of tuples, sizes of tuples
- Need to estimate the statistics of expression results
  - Can work as the input of another expression
  - To do so, we require additional statistics
    - E.g., the number of distinct values for an attribute
    - Selectivity of a predicate conditions

# How to Get Estimated Statistics

---

- Choice #1: Histograms
  - Maintain an occurrence count per value (or range of values) in a column
- Choice #2: Sketches
  - A probabilistic data structure that gives an approximate count for a given value
- Choice #3: Sampling
  - DMBS maintains a small subset of each table that it then uses to evaluate expressions to compute selectivity.
- Not covered in this lecture.
  - Let's assume we have a perfect estimator that can always return the actual number.

# Single-Relation Query Planning

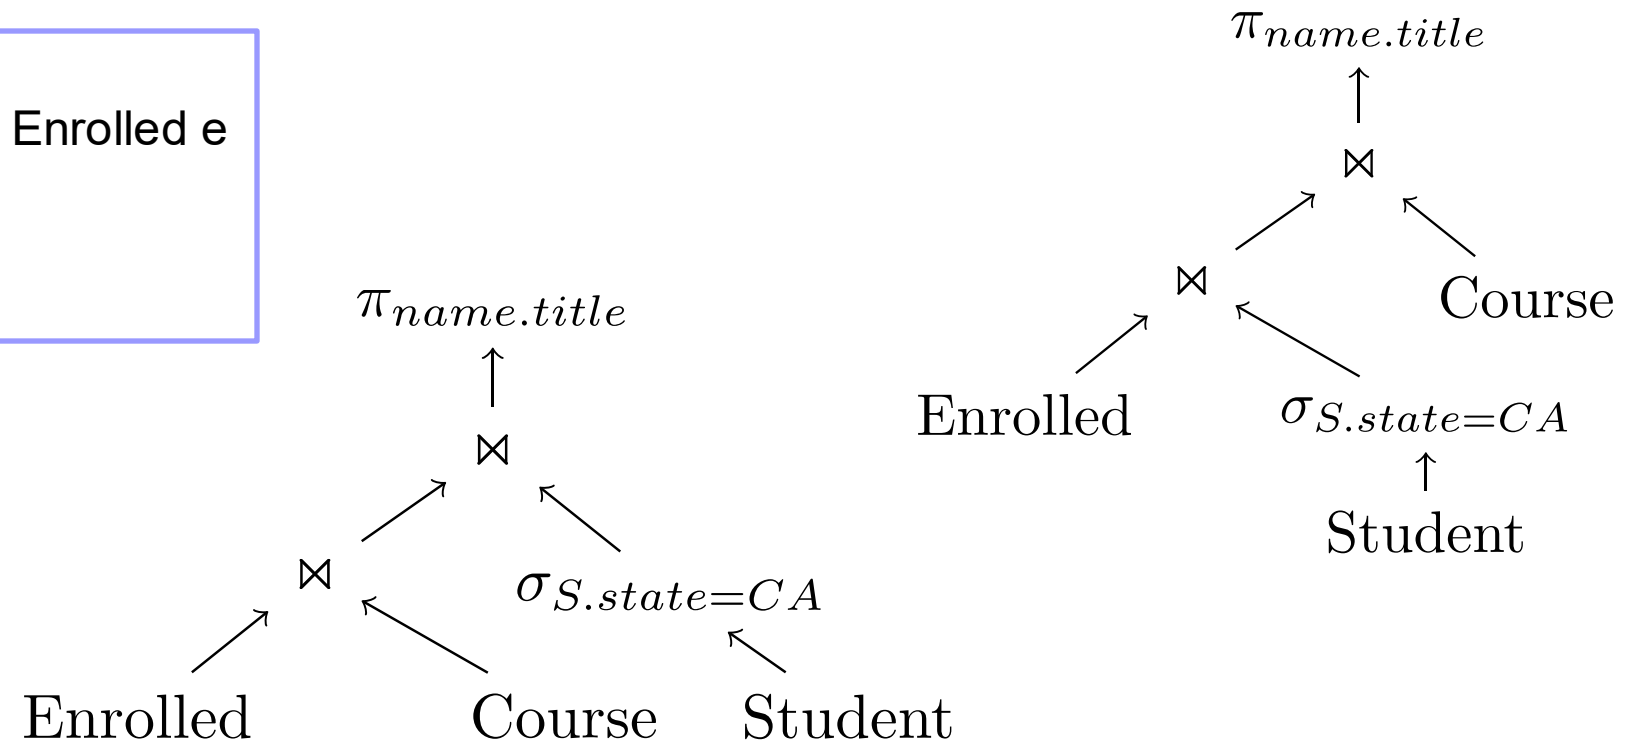
---

- Pick the best access method.
  - Sequential Scan  
e.g. , `Select * From R`, which requires accessing all records
  - Binary Search (clustered indexes)  
e.g. , Range filter conditions like `Select ... From R Where R.x <= 10;`
  - Index Scan  
e.g., Point filter conditions like `Select ... From R Where R.x = 'A';`
- Predicate evaluation ordering
  - Apply the predicates with indexes first to avoid a sequential scan
  - Apply the most restricted predicate first
- Simple heuristics are often good enough for this

# How to choose a better plan: Join Reordering

- Unlike predicate pushdown and projection pushdown, we cannot determine which relational expression is better after applying associative rules for multiple joins.


```
SELECT name, title
FROM Student s, Course c, Enrolled e
WHERE s.sid = e.sid
AND c.cid = e.cid
AND s.state = 'CA';
```



# Join Reordering


- Let's assume there are
  - 10000 records in the Enrolled relation
  - 50 records in the Course relation
  - 2000 records in the Student relation
  - Only 100 students are from CA
  - Every student enrolls in at most 10 courses

- Cost of the plan (The output size of each operation)
  - $Course \bowtie Enrolled$ : returns 10000 records.
  - The filter predicate returns 100 records.
  - The final join returns at most 1000 records.



# Join Reordering

- Let's assume there are
  - 10000 records in the Enrolled relation
  - 50 records in the Course relation
  - 2000 records in the Student relation
  - Only 100 students are from CA
  - Every student enrolls in at most 10 courses



- Cost of the plan (The output size of each operation)
  - The selective predicate returns 100 records.
  - The first join returns at most 1000 records.
  - The final join returns at most 1000 records.
- Assuming that generating one record requires a unit of time:
  - The first plan takes 11100 units
  - The second plan takes 2100 units

# Join Reordering

---

- Consider a chain join query:

$$R_1(x_1, x_2) \bowtie R_2(x_2, x_3) \bowtie \cdots \bowtie R_n(x_n, x_{n+1})$$

- There can be  $O(4^n)$  different join orders (Catalan number)
  - With 10 relations, total 4862 plans
  - With 20 relations, more than 1.7 billion plans

## Join Reordering (cont.)

- But there are a lot of duplicates for plans:

$$\left( \left( R_1(x_1, x_2) \bowtie R_2(x_2, x_3) \right) \bowtie R_3(x_3, x_4) \right) \bowtie \left( \left( R_4(x_4, x_5) \bowtie R_5(x_5, x_6) \right) \bowtie R_6(x_6, x_7) \right)$$

and

$$\left( \left( R_1(x_1, x_2) \bowtie R_2(x_2, x_3) \right) \bowtie R_3(x_3, x_4) \right) \bowtie \left( R_4(x_4, x_5) \bowtie \left( R_5(x_5, x_6) \bowtie R_6(x_6, x_7) \right) \right)$$

shares the same plan for evaluating the joins between  $R_1, R_2, R_3$

- The problem has **overlapping sub-problems** and show **optimal sub-structure**.

# Dynamic Programming!



# Dynamic Programming for Join Ordering

- Let  $cost[i, j]$  store the minimal cost for calculating chain query  $R_i \bowtie \dots \bowtie R_j$ , with  $plan[i, j]$  store the corresponding query plan. Assume the cost of calculating a join query is the size of the result.
  - When  $i > j$ , the problem is invalid
  - When  $i = j$ , return the relation  $R_i$  directly with the cost of  $|R_i|$
- When calculating the optimal plan for chain query  $R_i \bowtie \dots \bowtie R_j$ , we determine the position  $k$  for performing the last join
  - i.e., we calculate  $R_i \bowtie \dots \bowtie R_k$  and  $R_{k+1} \bowtie \dots \bowtie R_j$  first, and then calculate the join query
$$(R_i \bowtie \dots \bowtie R_k) \bowtie (R_{k+1} \bowtie \dots \bowtie R_j)$$
  - There are totally  $j - i$  different choices
- The cost of choosing  $k$  will be
$$cost[i, k] + cost[k + 1, j] + |R_i \bowtie \dots \bowtie R_j|$$

# Bottom-up Procedure

- To calculate the optimal cost for  $[i, j]$ , we first calculate all  $cost[l, m]$  with  $i \leq l \leq m \leq j$  and  $m - l < j - i$
- Then we try all possible  $k$  and keep only the optimal one.

```
Input:  $R_1, \dots, R_n$  in chain order;
for  $i \leftarrow 1$  to  $n$  do
    |  $cost[i, i] \leftarrow |R_i|$                                 // single relation cost
end
// Outer loop, set segment length
for  $L \leftarrow 2$  to  $n$  do
    // Middle loop, set the start index  $i$ 
    for  $i \leftarrow 1$  to  $n - L + 1$  do
        |  $j \leftarrow i + L - 1$ 
        |  $cost[i, j] \leftarrow \infty$ 
        // Inner loop, set the split point
        for  $k \leftarrow i$  to  $j - 1$  do
            |  $c \leftarrow cost[i, k] + cost[k + 1, j] + |R_i \bowtie \dots \bowtie R_j|$ 
            | if  $c < cost[i, j]$  then
            |     |  $cost[i, j] \leftarrow c;$ 
            |     |  $plan[i, j] \leftarrow k$ 
            | end
        end
    end
end
Output:  $cost[1, n]$  and query plan via  $plan$ 
```


# Complexity Analysis

---


- $O(n^2)$  memory cost
- $O(n^3)$  time complexity
  - When  $n = 20$ , the cost is 8000 instead of 1.7 billion.
- It is still costly if  $n$  is large.

# Left Deep Query Plans

- In left-deep query plans, the right-hand-side input for each join is a relation, not the result of an intermediate join




(a) Left-deep join tree



(b) Non-left-deep join tree

# Left Deep Query Plans (cont.)

- If only left deep query plans are considered, the number of query plans is significantly reduced.
  - For a chain query, the right-most relation must be  $R_1$  or  $R_n$
  - If we also fix the left-most relation to be  $R_1$ , the query plan is uniquely determined.
  - For  $n = 5$ , the plan is  $((R_1 \bowtie R_2) \bowtie R_3) \bowtie R_4 \bowtie R_5$



# Left Deep Query Plans (cont.)

- For calculating  $cost[i, j]$ , we only need to consider the right-most relation to be  $R_i$  or  $R_j$ 
  - No need to choose split point  $k$  anymore.
  - Reduce a factor of  $n$  for time complexity.

```
Input:  $R_1, \dots, R_n$  in chain order;
for  $i \leftarrow 1$  to  $n$  do
    |  $cost[i, i] \leftarrow |R_i|$                                 // single relation cost
end
// Outer loop, set segment length
for  $L \leftarrow 2$  to  $n$  do
    // Middle loop, set the start index  $i$ 
    for  $i \leftarrow 1$  to  $n - L + 1$  do
        |  $j \leftarrow i + L - 1$ 
        | // Choose  $R_i$  or  $R_j$  to be the right-most relation
        |  $c_1 \leftarrow cost[i, j - 1] + cost[j, j] + |R_i \bowtie \dots \bowtie R_j|$ 
        |  $c_2 \leftarrow cost[i, i] + cost[i + 1, j] + |R_i \bowtie \dots \bowtie R_j|$ 
        | if  $c_1 < c_2$  then
        |     |  $cost[i, j] \leftarrow c_1$ 
        |     |  $plan[i, j] \leftarrow j$ 
        | else
        |     |  $cost[i, j] \leftarrow c_2$ 
        |     |  $plan[i, j] \leftarrow i$ 
        | end
    end
end
Output:  $cost[1, n]$  and query plan via  $plan$ 
```

# Conclusion

---

- Query optimization is critical for a database system.
  - SQL -> Logical Plan -> Physical Plan
- The optimization step:
  - Write good SQL if possible.
  - Rule-based optimization for filtering logical plans.
    - Finding equivalent relational expressions
  - Cost-based optimization is used to select the best logical and physical plan.
    - A dynamic programming-based algorithm to avoid plan recomputation
- What is missing:
  - Some equivalent rules (read Database System Concepts, Section 13.2.1, and finish the practice exercises)
  - The cost estimation methods (Section 13.3)
- If you like this and want to make cash money in the database industry, consider earning a PhD in the database team at NTU.

## Reference

- [illegible]